Question #16   *Topic 1*

HOTSPOT -

You have two Azure Storage accounts named Storage1 and Storage2. Each account holds one container and has the hierarchical namespace enabled. The system has files that contain data stored in the Apache Parquet format.

You need to copy folders and files from Storage1 to Storage2 by using a Data Factory copy activity. The solution must meet the following requirements:

☞ No transformations must be performed.

☞ The original folder structure must be retained.

☞ Minimize time required to perform the copy activity.

How should you configure the copy activity? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Source dataset type:

```
Binary
Parquet
Delimited text
```

Copy activity copy behavior:

```
FlattenHierarchy
MergeFiles
PreserveHierarchy
```

Question #17   *Topic 1*

You have an Azure Data Lake Storage Gen2 container that contains 100 TB of data.

You need to ensure that the data in the container is available for read workloads in a secondary region if an outage occurs in the primary region.

The solution must minimize costs.

Which type of data redundancy should you use?

   A. geo-redundant storage (GRS)

   B. read-access geo-redundant storage (RA-GRS)

   C. zone-redundant storage (ZRS)

   D. locally-redundant storage (LRS)

You plan to implement an Azure Data Lake Gen 2 storage account.

You need to ensure that the data lake will remain available if a data center fails in the primary Azure region. The solution must minimize costs. Which type of replication should you use for the storage account?

    A. geo-redundant storage (GRS)

    B. geo-zone-redundant storage (GZRS)

    C. locally-redundant storage (LRS)

    D. zone-redundant storage (ZRS)

HOTSPOT -

You have a SQL pool in Azure Synapse.

You plan to load data from Azure Blob storage to a staging table. Approximately 1 million rows of data will be loaded daily. The table will be truncated before each daily load.

You need to create the staging table. The solution must minimize how long it takes to load the data to the staging table.

How should you configure the table? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Distribution:
- Hash
- Replicated
- Round-robin

Indexing:
- Clustered
- Clustered columnstore
- Heap

Partitioning:
- Date
- None